

Short Term Interest Rate Innovations Across Monetary Regimes

Diogo Ribeiro^{1,2,*} ORCID: 0009-0001-2022-7072

Department of Data Science and Research, MySense, London, UNITED KINGDOM.

School of Media Arts and Design, Polytechnic of Porto
PORTUGAL.

**Corresponding Author: dfr@esmad.ipp.pt*

Draft

Status: Preliminary research design. This draft records the research question, replication design, and planned extensions. It does not report empirical estimates, risk prices, standard errors, or pricing-error magnitudes.

Abstract

This paper studies whether short-term interest-rate innovations help explain the cross section of stock-market anomaly returns, and whether that relation is stable across later monetary regimes. The design begins with a strict replication of Maio and Santa-Clara (2017) and separates exact replication, documented reconstruction, and new extension evidence. The planned extensions test temporal stability, regime dependence, weak-factor sensitivity, and comparator-model performance. The current draft is preliminary: the publication targets and source-access requirements are registered, but source-compatible empirical panels and generated result tables are not yet available. Numerical findings are therefore withheld until the baseline and extension artifacts can be generated and audited.

1 Introduction

Short-term interest rates are plausible state variables in intertemporal asset-pricing models because they summarize investment opportunities that matter for both discount rates and hedging demands. If portfolios formed on accounting or return-based anomalies have different exposures to those state variables, a rate innovation may help organize average excess returns beyond the market factor. The economic question is not only whether such a factor prices anomaly returns in one sample, but whether the relation survives changes in monetary regimes and source definitions.

Maio and Santa-Clara (2017) provide the starting point. Their article studies short-term interest-rate innovations and stock-market anomalies, and this paper uses that design as the baseline object to be reproduced before any extension is interpreted. The strict baseline window is January 1972 through December 2013. The extension window begins in January 2014 and is evaluated separately from the replication verdict.

The paper has three aims. First, it rebuilds the original baseline in a way that distinguishes exact replication from documented reconstruction. Second, it asks whether the short-rate factor is stable after the publication sample and across monetary regimes. Third, it evaluates whether any short-rate pricing evidence is robust to weak-factor diagnostics and comparator-model discipline.

The high-frequency policy-information decomposition is kept as an appendix design until usable event data and monthly component artifacts exist.

The current evidence status is intentionally limited. The article and supplement targets have been registered, and the software pipeline contains data, estimation, diagnostic, manuscript, and release gates. The empirical inputs needed to populate the baseline tables are still incomplete, so the results sections below state pending-result conditions rather than estimates.

2 Related literature

The first literature stream motivates the state variable. The ICAPM shows why assets can earn premia for covariance with variables that forecast future investment opportunities (Merton, 1973). Interest-rate applications of this idea are directly relevant because Lioui and Maio (2014) model interest-rate risk as a priced source of cross-sectional equity returns, and Maio and Santa-Clara (2017) apply short-rate innovations to anomaly portfolios.

The second stream concerns anomaly pricing and model comparison. The baseline article compares a short-rate ICAPM with standard factor models, so the extension must be evaluated against asset-pricing benchmarks rather than against the CAPM alone. Fama and French (2015) add profitability and investment factors to the Fama-French framework, while Hou et al. (2015) propose an investment approach that summarizes many anomalies with market, size, investment, and profitability factors. These benchmarks make the extension question sharper: does the short-rate factor add robust pricing content once comparator models, asset intersections, and portfolio definitions are held fixed?

The third stream concerns inference. Two-pass cross-sectional tests are exposed to generated-regressor and weak-factor problems. Fama and MacBeth (1973) provide the classic cross-sectional testing framework, Shanken (1992) derives corrections for beta-estimation error, and Kan and Zhang (1999) show that useless factors can appear priced in two-pass tests. Lewellen et al. (2010) argue that high cross-sectional fit and small pricing errors can be misleading when test assets share a strong characteristic structure. This paper therefore treats beta dispersion, rank, spanning, and influence diagnostics as conditions for interpretation rather than optional robustness checks.

The fourth stream concerns monetary regimes and high-frequency identification. Bernanke and Kuttner (2005) measure equity-market responses to unanticipated Federal Reserve policy actions, while Jarociński and Karadi (2020) separate policy and central-bank-information components using high-frequency announcement comovement. Swanson and Williams (2014) show why lower-bound episodes can change the informativeness of short-term rates for the broader term structure. Those papers motivate separating the aggregate AR rate innovation from identified announcement components and from regime-specific pricing evidence. Structural-change methods such as Bai and Perron (1998) provide a way to distinguish predeclared regime boundaries from estimated breaks.

The final stream is data vintage discipline. Official factor archives document that published factor and portfolio series can change as databases and construction rules are revised (French, 2026). This paper therefore adds value beyond the original article by freezing a table-level replication target before asking whether the short-rate evidence generalizes over time and across regimes.

3 Original framework and economic mechanism

The original framework is a discrete-time ICAPM application in which the market excess return is paired with an innovation in a short-term interest-rate state variable. The article verifies two state

variables: the federal funds rate and the three-month Treasury-bill rate. It motivates both series as proxies for investment opportunities and constructs innovations from first-order autoregressions.

For the federal-funds version, the article reports the state-variable law of motion as

$$FFR_{t+1} = 0.000 + 0.991FFR_t + u_{t+1}^{FFR}, \quad R^2 = 0.98,$$

which this manuscript indexes for the return month as

$$u_t^{FFR} = FFR_t - \hat{a} - \hat{\rho}FFR_{t-1}.$$

This convention fixes the information timing: the return regression for month t uses the innovation from the same realized month and never relies on an implicit shift. The article reports the analogous Treasury-bill equation as

$$TB_{t+1} = 0.000 + 0.992TB_t + u_{t+1}^{TB}, \quad R^2 = 0.98.$$

For test asset i , excess return $R_{i,t}^e$, market factor $R_{M,t}$, and federal-funds innovation u_t^{FFR} , the verified first-pass time-series regression is

$$R_{i,t}^e = \delta_i + \beta_{i,M}R_{M,t} + \beta_{i,FFR}u_t^{FFR} + \varepsilon_{i,t}.$$

The benchmark second-pass regression is a single OLS cross-sectional regression of full-sample average excess portfolio returns on full-sample first-pass betas, without an intercept:

$$\bar{R}_i^e = \hat{\beta}_{i,M}\lambda_M + \hat{\beta}_{i,FFR}\lambda_{FFR} + \alpha_i^{CS}.$$

Here δ_i is the time-series intercept, α_i^{CS} is the cross-sectional pricing error, and λ_{FFR} is the federal-funds innovation risk price. This benchmark is not described as repeated monthly Fama–MacBeth estimation unless a separate supplement target or robustness artifact explicitly implements repeated cross sections. The no-intercept restriction is part of the article’s benchmark economic test, not a default regression convention.

The economic mechanism is an ICAPM-style hedging argument. A rate innovation can matter if it captures news about future investment opportunities and if anomaly portfolios differ in their exposure to that news. That statement is weaker than a monetary-policy interpretation. An autoregressive residual may combine policy news, information about the macroeconomy, data revisions, and persistence in the short-rate series. The paper therefore treats the baseline factor as a rate innovation until a separate high-frequency design justifies narrower language.

The article’s primary test assets are value-weighted decile portfolios sorted on book-to-market, equity duration, earnings-to-price, long-term reversal, investment-to-assets, property-plant-and-equipment investment, and inventory growth. It also studies the joint system of all seven anomaly families and uses equal-weighted portfolios and additional appendix designs as robustness checks. The benchmark inference reports Shanken-corrected risk-price statistics, bootstrap p-values, a joint pricing-error statistic, and cross-sectional fit metrics.

The mechanism is not one channel. Short-rate innovations may proxy for discount-rate news, cash-flow news, financing-cost changes, working-capital conditions, monetary-policy information, or intertemporal hedging demand. The equity-duration channel is the clearest family-level mechanism because a rate innovation changes the discounting of cash flows at different horizons. For the other anomaly families, the paper predeclares exposure comparisons but does not force a directional sign unless the article-level result and generated portfolio definition support it.

The meaning of a numerical rate innovation can also change across regimes. Conventional positive-rate policy, lower-bound periods, asset-purchase and forward-guidance episodes, and

Table 1: Replication target and economic hypotheses

Claim	Role	Primary test	Interpretation rule
R1a–R1f	Replication targets	Layer-specific article statistics within tolerances	Numerical recovery only
H1	Incremental pricing content	Improvement versus strongest comparator	Economic thresholds required
H2	Temporal stability	Post-publication compatibility	Sign and magnitude matter
H3	Regime stability	Separate beta, fitted-premium, and error tests	Non-rejection is insufficient
H4	Factor-strength condition	Beta dispersion and rank diagnostics	Weak factors limit interpretation

Source: Authors’ replication protocol.

inflation-driven tightening can attach different information to the same measured short-rate movement. The baseline rate residual therefore remains an aggregate state-variable innovation; policy and information language is reserved for the optional announcement-based appendix design.

Some implementation details remain ambiguous at the source-file level. The article identifies the Federal Reserve Bank of St. Louis, Kenneth French’s data library, Robert Stambaugh’s web page, and Lu Zhang as sources, but the exact download files and archive vintages still have to be frozen before the manuscript can claim exact computational replication.

4 Data and replication design

The data design separates exact inputs, approximate reconstructions, and extension inputs. Exact inputs require a verified article definition, source, vintage, transformation, and estimator role. Approximate reconstructions are documented public or author-data substitutes and are never labelled as exact replication. Extension inputs support post-publication, regime, and high-frequency analyses after their source terms and transformations are frozen.

The baseline sample runs from 1972-01 through 2013-12. The temporal extension begins in 2014-01, with revised-history checks kept separate from the baseline verdict. Every source must record access status, redistribution status, and its role in strict replication before it can support a claim. Exact unavailable inputs lead to a missing-input classification rather than an empirical contradiction.

The short-rate registry fixes the documented reconstruction series before the AR model is estimated. The federal-funds reconstruction uses FRED series FEDFUNDS, an effective federal funds rate, monthly provider average, annualized percent units, no seasonal adjustment, no imputation, and a vintage audit at retrieval. The Treasury-bill reconstruction uses TB3MS as the monthly secondary-market candidate and treats DTB3 only as a sensitivity candidate unless the article or supplement verifies a daily aggregation rule.

Portfolio continuity is the largest post-publication constraint. For all seven registered anomaly families, the extension path is full security-level reconstruction using CRSP returns, delisting returns, Compustat accounting inputs, original formation dates, breakpoints, value-weighting rules, and holding-period rules. Current public portfolio catalogues may be reported as documented substitutes, but they are excluded from exact-replication claims unless source compatibility is independently verified.

Short-rate series are stored in their source units and transformed only in named innovation

Table 2: Data sources and sample-construction status

Variable or family	Source	Frequency	Transformation	Status
Article and supplement	Publisher files	Document	Target extraction	Exact local access
Federal funds rate	Federal Reserve Bank of St. Louis	Monthly	First-order innovation	Source located, series pending
Treasury-bill rate	Federal Reserve Bank of St. Louis or author source	Monthly	First-order innovation	Series ambiguous
Market and risk-free returns	Kenneth French Data Library	Monthly	Excess-return factors	Source located, vintage pending
Anomaly decile portfolios	Lu Zhang and cited sources	Monthly	Excess returns after alignment	Files pending
Comparator factors	French, Stambaugh, and q-factor sources	Monthly	Common-intersection factors	Source located, vintage pending
Announcement components	Documented research dataset	Event	Monthly aggregation	Extension input pending

Source: Authors' replication protocol.

steps. Market and portfolio returns are aligned to the canonical monthly panel before excess returns, factor construction, or model comparison. Missing returns and shock observations are not forward filled. Comparator models use identical asset-date intersections within each comparison, and delisting or survivorship treatment is recorded from verified source documentation rather than inferred from the data file alone.




Figure 1: Data flow from sources to manuscript evidence

Source: Authors' replication protocol.

The replication target is table-level. Each published table or appendix table is mapped to a source location, model, estimator, tolerance rule, and status. A generated statistic can be labelled reproduced only after the same definition and estimator are linked to an auditable output. Nearby public data may be useful for documented reconstruction, but they cannot change the strict replication label. The replication target is decomposed into short-rate innovations, first-pass betas, cross-sectional risk prices, pricing errors and fit, comparator-model results, and supplementary robustness targets.

5 Econometric methods

The baseline empirical design starts with the article’s first-order short-rate innovation equations. Alternative autoregressive orders, alternative rate definitions, local-level filters, and announcement-surprise measures are robustness or extension specifications, not replacements for the strict baseline.

First-pass time-series regressions estimate portfolio exposures to the market excess return and the selected short-rate innovation. The notation separates the time-series intercept δ_i , factor loadings β_i , residuals $\varepsilon_{i,t}$, cross-sectional pricing errors α_i^{CS} , and risk prices λ . The second pass is the article’s verified no-intercept OLS regression of average returns on estimated betas unless a supplement target explicitly calls for an unrestricted zero-beta-rate variant, GLS estimator, or repeated monthly cross-sectional estimator. Shanken corrections apply to the reported full-sample risk-price estimator in the benchmark design; bootstrap procedures are labelled by resampling unit before any p-value is interpreted.

Model comparison reports pricing-error loss and fit on common asset-date intersections. The registered outcomes are root mean squared pricing error, mean absolute pricing error, maximum absolute pricing error, cross-sectional fit, and joint pricing-error tests when the estimator supports them. Statistical significance is not enough: every supported result must also clear the registered economic-magnitude criterion. For H1, materiality is defined against the strongest registered non-short-rate comparator by baseline RMSE on the common intersection. The short-rate model must reduce RMSE and MAE by at least 10 percent and reduce the maximum absolute pricing error by at least 0.25 monthly percentage points relative to that comparator.

Weak-factor diagnostics evaluate whether short-rate betas contain enough cross-sectional variation to support a pricing interpretation. The planned diagnostics include beta dispersion, matrix rank, singular values, condition measures, factor-spanning checks, influence diagnostics, and leave-one-family systems. A short-rate risk price that fails these diagnostics is interpreted as weakly identified even when fit metrics appear favorable. The factor fails the confirmatory strength gate if the beta matrix loses rank, the smallest singular value is below 10 percent of the largest singular value, short-rate beta dispersion is below 10 percent of market-beta dispersion, the condition number exceeds 30, leave-one-family systems reverse the fitted-premium spread or materiality classification, or robust intervals span economically positive and negative fitted-premium effects.

Temporal and regime methods are designed as extensions. The post-publication analysis has frozen-parameter and refitted variants, with no-lookahead rules for factor construction, portfolio membership, model vintage, and sample intersections. Regime analysis separates four questions: stability of the rate innovation construction, stability of portfolio betas, stability of cross-sectional risk prices, and stability of pricing errors. Pooled time-series interactions test beta stability; regime-specific or locally estimated second-pass models test risk-price stability. Stability is not inferred from a failure to reject equality.

Regime risk-price stability is evaluated through the rate-attributable fitted premium $\pi_{i,s} = \beta_{i,FFR,s} \lambda_{FFR,s}$, not through an absolute bound on $\lambda_{FFR,s}$. This object is invariant to rescaling the short-rate innovation. The predeclared equivalence bounds require fitted-premium changes no larger than 0.25 monthly percentage points, fitted-premium dispersion changes no larger than 25 percent, RMSE deterioration no larger than 10 percent, maximum pricing-error deterioration no larger than 0.25 monthly percentage points, and selected-fit deterioration no larger than 0.10 under the fit contract.

Cross-sectional fit is mechanical rather than discretionary. The primary fit statistic is the article’s exact definition if verified; otherwise the fallback is a no-intercept pricing-error pseudo- R^2 , $1 - \sum_i (\alpha_i^{CS})^2 / \sum_i (\bar{R}_i^e)^2$, reported only when the denominator is positive and all models use the same asset-date intersection.

Inference uses Shanken-adjusted covariance for the benchmark full-sample second-pass risk-price estimator. Bootstrap p-values use moving blocks of monthly observations, 10,000 repetitions, and either the configured serial-dependence block rule or a labelled 12-month fallback. Equivalence claims use 95 percent confidence intervals that must lie wholly inside the declared bounds.

The primary regime registry contains five regimes: conventional positive-rate policy, effective lower bound and first QE, first normalisation, pandemic lower-bound and QE, and inflation-driven tightening. Transition months are assigned to the new regime, boundary-shift checks are plus or minus three months, and short-sample flags are retained for the two post-pandemic regimes.

The high-frequency decomposition is optional rather than part of the main contribution. If included, it starts from high-frequency components constructed under an explicit announcement-based identification design after source terms and component definitions are frozen. The monthly aggregation can create a sparse factor with many no-event months, so the component factors must pass their own weak-factor diagnostics before entering the main argument.

Out-of-sample evaluation defines the forecast origin, training window, refit schedule, benchmark models, loss functions, and interpretation before results are inspected. It is an appendix robustness exercise unless it produces evidence strong enough to motivate a separate contribution. Multiplicity adjustments are predeclared within replication, baseline pricing, temporal extension, regime stability, and weak-factor families.

Table 3: Method-to-claim map

Claim	Primary method	Decision rule
R1a–R1f	Layer-specific comparison to registered article statistics	Recovery within declared tolerances
H1	No-intercept two-pass pricing and model comparison	Must beat strongest comparator by materiality thresholds
H2	Frozen and refitted post-publication evaluation	Sign, magnitude, and loss compatibility must hold
H3	Separate beta, fitted-premium, and pricing-error stability tests	Estimates must remain inside numerical bounds
H4	Beta dispersion, rank, and robust-inference diagnostics	Strength gates must pass before interpretation

Source: Authors’ replication protocol.

6 Planned baseline replication

Baseline empirical results are pending. The current replication audit is blocked by missing source-compatible inputs and generated baseline artifacts. The planned baseline analysis uses the strict monthly baseline, the registered baseline and comparator system, table-level generated-cell audits, and the registered anomaly portfolios. Uncertainty and economic magnitude remain unreported until generated standard errors and replicated table cells exist.

This pending status does not evaluate the published evidence. It states only that the repository cannot yet generate the necessary baseline tables from source-compatible inputs. Once the missing inputs are registered, the same audit will classify each target as reproduced, approximately reproduced, contradicted, not attempted, or still blocked by missing input. A completed attempt that cannot recover a published target is the only setting in which the manuscript will use a non-reproducibility label.

7 Planned weak-factor and model-comparison diagnostics

The diagnostic section will report whether any baseline pricing result survives alternative rate innovations, covariance estimators, portfolio weighting, influence checks, weak-factor diagnostics, and comparator-model tests. The first question is mechanical: do estimated short-rate betas vary enough across assets for the second pass to identify a risk price? The second question is economic: does the factor reduce pricing errors by a material amount relative to registered comparator models?

Robustness results are pending because baseline estimates are pending. The planned diagnostics use the generated strict baseline, the registered robustness family, diagnostic comparisons over baseline pricing artifacts, and the replication-audit assets. Uncertainty is governed by weak-factor and familywise diagnostics, and economic magnitude is blocked until baseline estimates exist.

8 Planned post-publication evidence

The temporal extension asks whether the baseline relation is stable after the publication sample. The baseline endpoint remains 2013-12, and the extension starts in 2014-01. The design prevents revised data or later observations from changing the baseline replication label.

Post-publication results are pending. The planned analysis uses the extension-compatible monthly panel, frozen and refitted versions of the registered two-factor model, source-compatible anomaly portfolios, and predeclared temporal loss comparisons. Economic interpretation requires sign compatibility, risk-price magnitude compatibility, and pricing-error performance relative to the frozen baseline and comparator models.

9 Planned monetary-regime stability

The regime extension asks whether the short-rate pricing relation changes across predeclared monetary regimes. It separates stability of factor construction, portfolio betas, risk prices, and pricing errors. Pooled time-series interactions are evidence about exposure stability, not by themselves evidence about cross-sectional risk-price stability. Risk-price stability requires regime-specific or locally estimated second-pass models.

Regime results are pending. The planned analysis uses the extension-compatible monthly panel, predeclared regime indicators, source-compatible anomaly portfolios, interaction tests for betas, separate second-pass estimates for risk prices, pricing-error comparisons, Holm-adjusted families, equivalence classifications, boundary-shift checks, and exploratory break diagnostics.

10 Discussion and limitations

The main limitation is input availability. Exact rate definitions, portfolio files, generated factor panels, baseline estimator outputs, extension panels, and event-level shock data are not yet complete. This limitation is a research status condition, not evidence for or against the baseline article.

The second limitation is interpretive. Even with a replicated risk price, a priced covariance with a rate innovation would not by itself establish a narrow monetary-policy mechanism. The design therefore separates pricing evidence from identification evidence and treats policy-information decomposition as an appendix design until event data exist.

The third limitation is econometric. Cross-sectional pricing is fragile when betas have little dispersion, samples are short within regimes, or comparator models are evaluated on different

asset-date intersections. The paper’s diagnostic structure is intended to make those limits visible before substantive interpretation.

11 Conclusion

This preliminary draft converts the repository protocol into an academic paper structure while preserving the evidence discipline required for replication. The paper will first ask whether the published short-rate innovation pricing evidence can be reproduced from source-compatible inputs. It will then ask whether the relation is stable after the publication sample and across monetary regimes, and whether any pricing relation survives weak-factor and model-comparison diagnostics. No empirical conclusion is reported until the registered inputs and generated artifacts are available.

A Table-level replication audit

The table-level audit maps the article and supplement targets to replication statuses. The current audit covers the article tables and supplement appendix tables through Table A.14, and every target is blocked by missing generated baseline inputs. The generated audit records the current missing-input classification for each target.

Table 4: Original table groups and replication targets

Original evidence group			Registered targets	Current classification
Article descriptive and pricing tables			TBL_001–TBL_009	Missing generated inputs
Supplement interest-rate ro-bustness			APP_TBL_A01– APP_TBL_A06	Missing generated inputs
Supplement portfolio and model extensions			APP_TBL_A07– APP_TBL_A14	Missing generated inputs

Source: `research/table_target_manifest.csv`.

B Data access and reproducibility

Restricted article and supplement files remain local-only. The committed repository records meta-data and source-access status, but it does not redistribute copyrighted source documents, prompt files, licensed data extracts, local catalogs, temporary artifacts, or machine-specific environment manifests. The source-only release gate reports zero critical restricted-path issues and one major missing-input issue.

Source-access requirements, release artifacts, and validation status are recorded outside the main argument in the repository documentation and generated release manifests.

C Hypothesis registry

The confirmatory registry contains a replication target and hypotheses for incremental pricing content, post-publication compatibility, regime equivalence, and weak-factor strength. The registry distinguishes null, alternative, primary outcome, estimator, sample, test assets, economic direction

or equivalence criterion, multiplicity family, and supported, unsupported, or inconclusive interpretation. Exploratory structural-break alignment is recorded separately from confirmatory claims.

D Table and figure architecture

The manuscript table and figure map defines the planned main-text evidence sequence before result writing begins. The architecture prioritizes a source-status table, baseline diagnostic table, replication-comparison table, weak-factor diagnostic table, baseline pricing table, comparator table, temporal-extension table, and regime-stability table. The planned main figures cover rate innovations and regimes, beta-return relations, rolling or expanding risk prices, pricing-error comparisons, and regime equivalence intervals.

The LaTeX shells and generation contracts contain no empirical cells and cannot be mistaken for output. Every shell is linked to a machine-readable source and an evidence gate. Result writing remains blocked until the required generated artifact exists.

E Detailed out-of-sample design

The out-of-sample design freezes an initial training endpoint at 1999-12 and uses a predeclared annual refit schedule. Forecast records must store model vintage, training dates, factor definition, asset universe, benchmark model, and loss function. Results from this design will be interpreted as evidence about stability of pricing errors, not as a substitute for the baseline replication audit.

F Optional policy-information decomposition

The policy-information decomposition is not part of the main contribution until legally usable event data, component labels, monthly aggregation artifacts, and component-level weak-factor diagnostics are available. If included, the components are high-frequency components constructed under an explicit announcement-based identification design, not automatically identified objects. The monthly aggregation can create sparse factors with many no-event months, so the decomposition must pass factor-strength diagnostics before entering the main pricing argument.

G Additional diagnostics

Additional diagnostics include beta dispersion, singular values, rank checks, factor spanning, influence diagnostics, leave-one-family systems, boundary-shift checks, and multiple-testing adjustments. Diagnostic outputs remain pending until the relevant baseline and extension artifacts exist.

References

- Bai, J. and Perron, P. (1998). Estimating and testing linear models with multiple structural changes. *Econometrica*, 66(1):47–78.
- Bernanke, B. S. and Kuttner, K. N. (2005). What explains the stock market’s reaction to federal reserve policy. *Journal of Finance*, 60(3):1221–1257.
- Fama, E. F. and French, K. R. (2015). A five-factor asset pricing model. *Journal of Financial Economics*, 116(1):1–22.

- Fama, E. F. and MacBeth, J. D. (1973). Risk, return, and equilibrium: Empirical tests. *Journal of Political Economy*, 81(3):607–636.
- French, K. R. (2026). Kenneth r. french data library. https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/data_library.html. Accessed 2026-07-31.
- Hou, K., Xue, C., and Zhang, L. (2015). Digesting anomalies: An investment approach. *The Review of Financial Studies*, 28(3):650–705.
- Jarociński, M. and Karadi, P. (2020). Deconstructing monetary policy surprises: The role of information shocks. *American Economic Journal: Macroeconomics*, 12(2):1–43.
- Kan, R. and Zhang, C. (1999). Two-pass tests of asset pricing models with useless factors. *Journal of Finance*, 54(1):203–235.
- Lewellen, J., Nagel, S., and Shanken, J. (2010). A skeptical appraisal of asset-pricing tests. *Journal of Financial Economics*, 96(2):175–194.
- Lioui, A. and Maio, P. (2014). Interest rate risk and the cross section of stock returns. *Journal of Financial and Quantitative Analysis*, 49(2):483–511.
- Maio, P. F. and Santa-Clara, P. (2017). Short-term interest rates and stock market anomalies. *Journal of Financial and Quantitative Analysis*, 52(3):927–961.
- Merton, R. C. (1973). An intertemporal capital asset pricing model. *Econometrica*, 41(5):867–887.
- Shanken, J. (1992). On the estimation of beta-pricing models. *The Review of Financial Studies*, 5(1):1–33.
- Swanson, E. T. and Williams, J. C. (2014). Measuring the effect of the zero lower bound on medium- and longer-term interest rates. *American Economic Review*, 104(10):3154–3185.